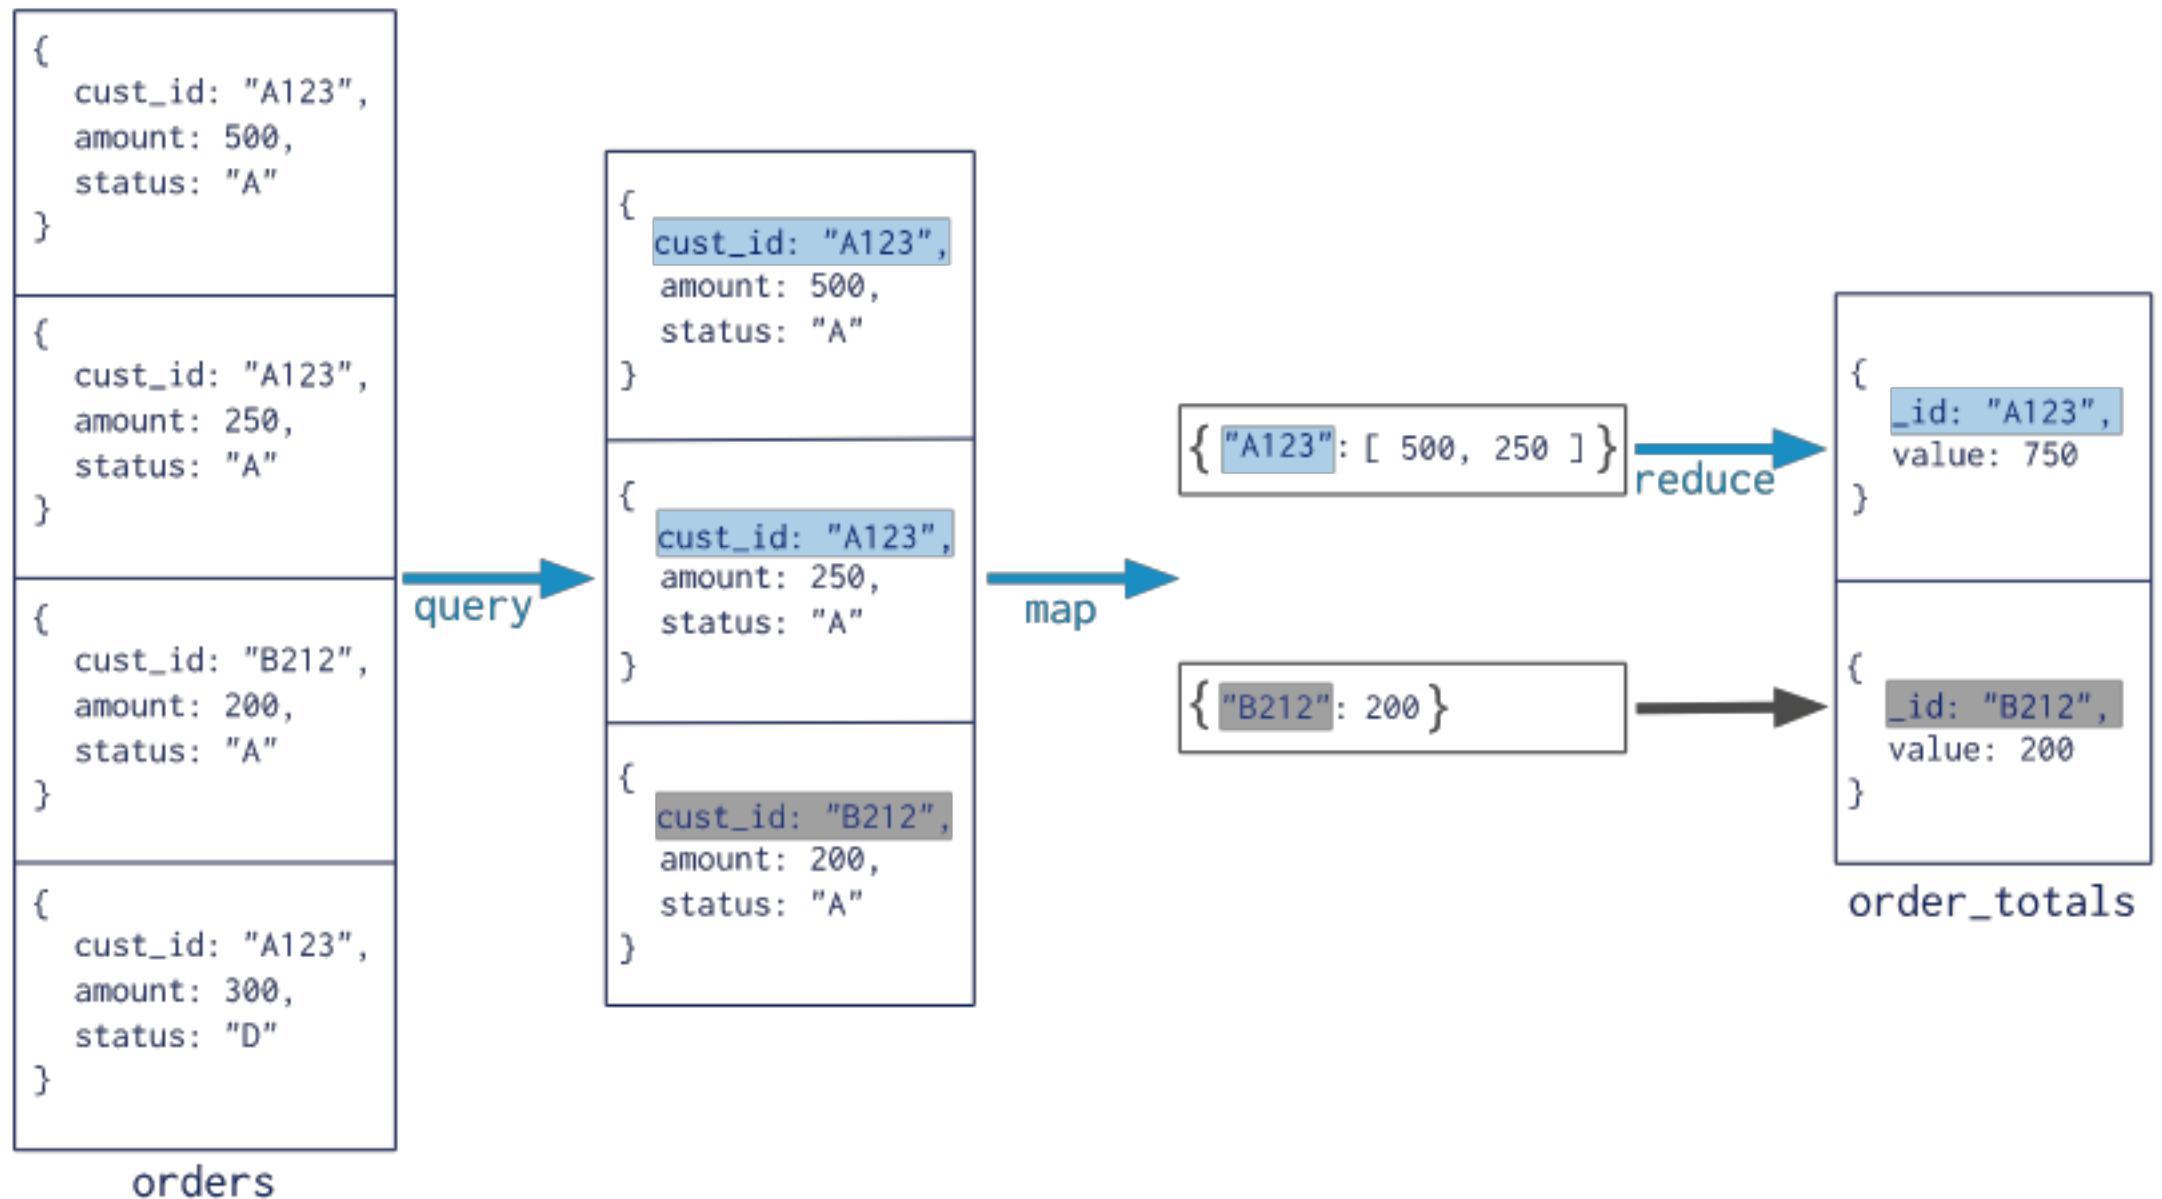


Collection
↓
db.orders.mapReduce(
 map → function() { emit(this.cust_id, this.amount); },
 reduce → function(key, values) { return Array.sum(values) },
 query → {
 output → query: { status: "A" },
 out: "order_totals"
 }
)



This map-reduce operation applies the **map** phase to documents in the collection that match some query condition (i.e., customer IDs).

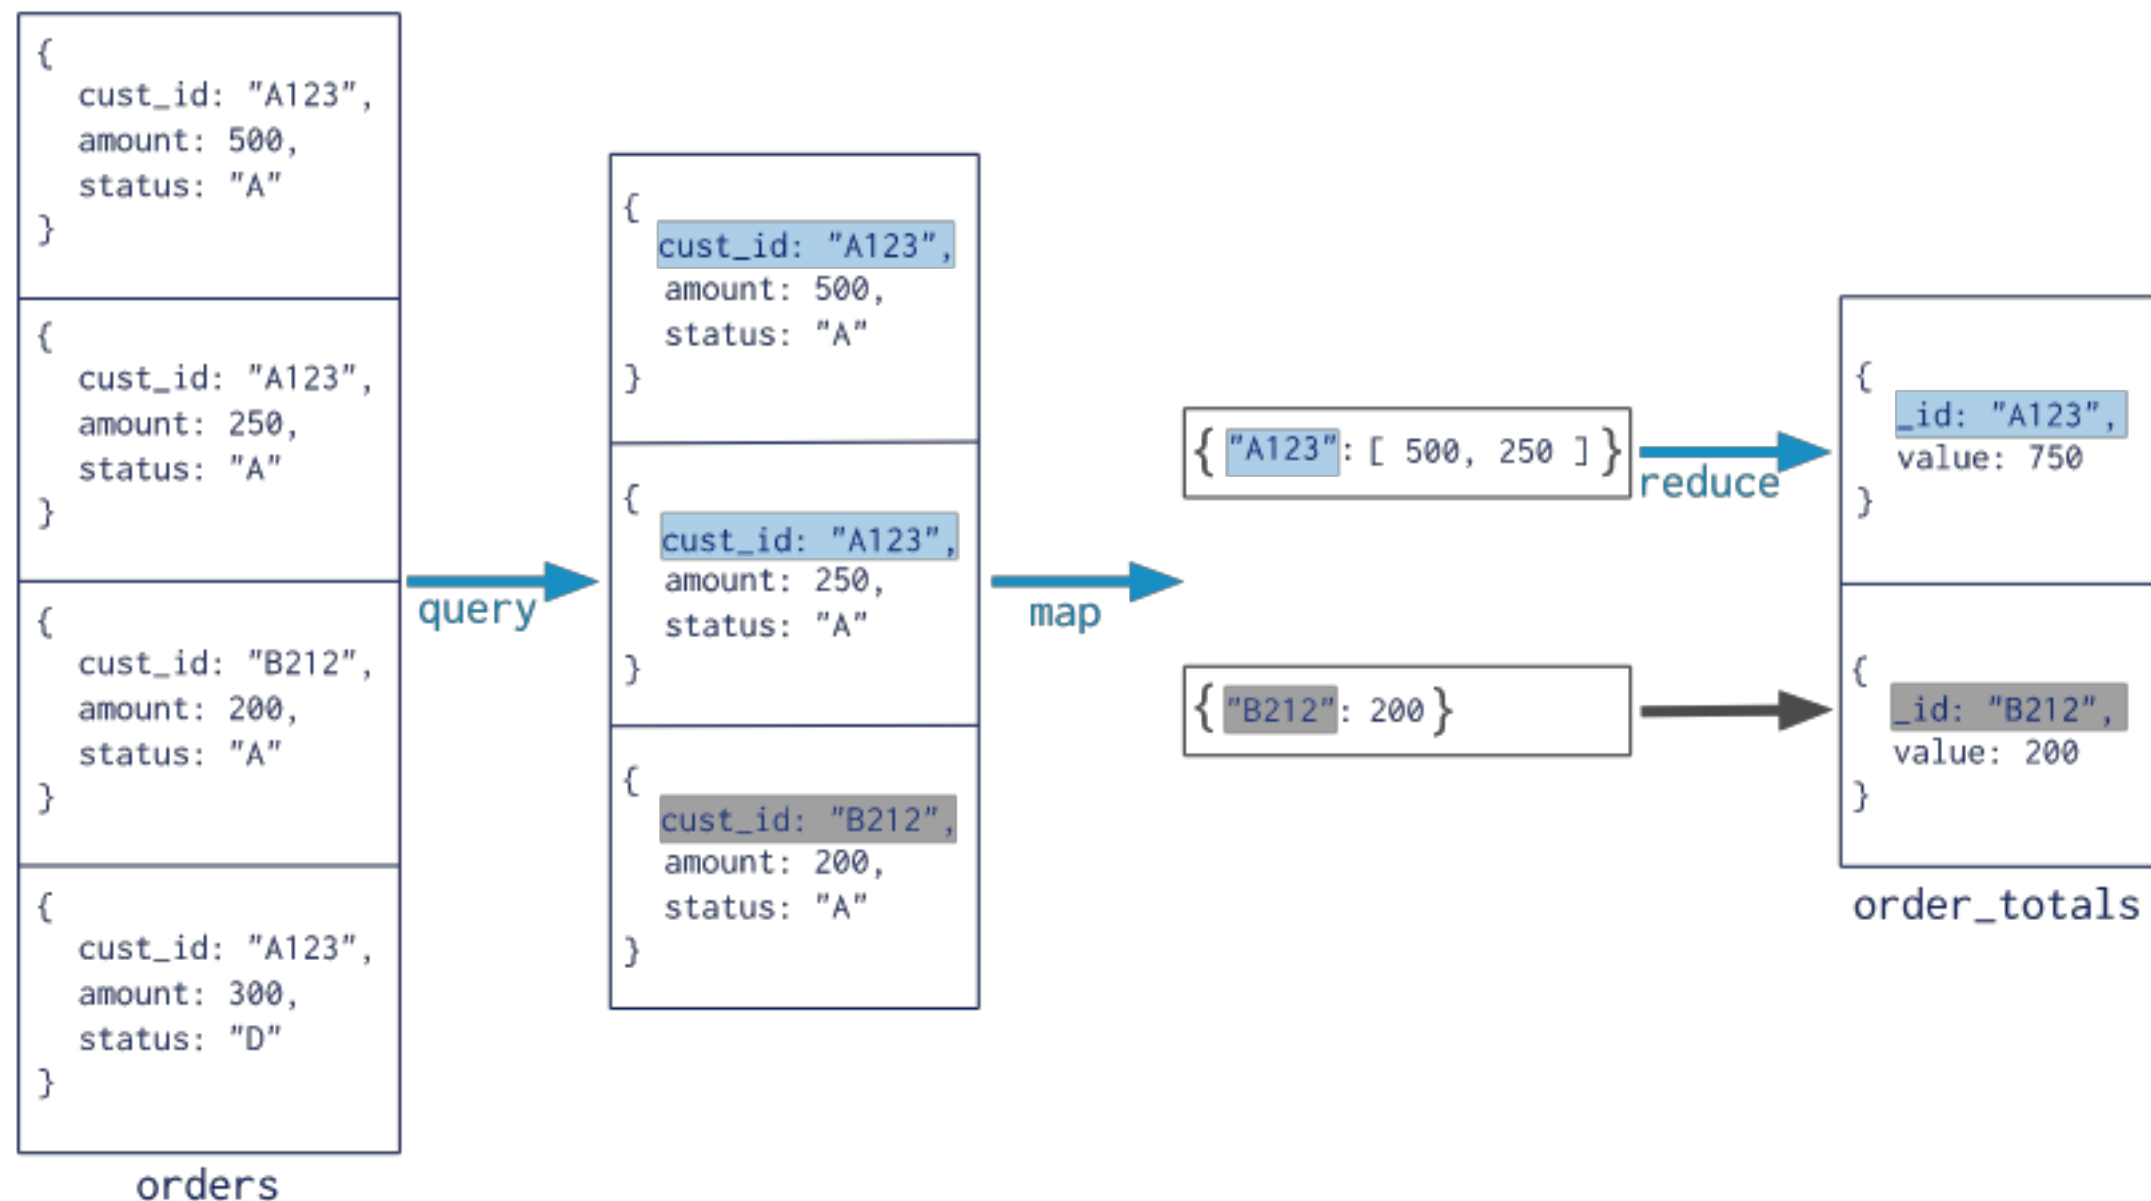
For those keys (customer IDs) with multiple values (i.e., orders), the **reduce** phase collects and condenses the aggregated data.

Finally, the results are aggregated and stored in some collection.

This map-reduce operation applies the **map** phase to documents in the collection that match some query condition (i.e., customer IDs).

For those keys (customer IDs) with multiple values (i.e., orders), the **reduce** phase collects and condenses the aggregated data.

Finally, the results are aggregated and stored in some collection.



BY THE END OF THIS SESSION YOU SHOULD BE ABLE TO:

- Conceptual:
 - Identify the 4Vs of Big Data, defining each criterion.
 - Understand the difference between structured and unstructured data.
 - Differentiate between serial and parallel processing.
 - Explain parallelization in terms of its advantages in dealing with Big Data.
 - Identify the map and reduce phases in MapReduce and explain what is happening in both parts.

